

Problem 2: CART-RF-ANN

An Insurance firm providing tour insurance is facing higher claim frequency. The management decides to collect data from the past few years. You are assigned the task to make a model which predicts the claim status and provide recommendations to management. Use CART, RF & ANN and compare the models' performances in train and test sets.

2.1 Data Ingestion: Read the dataset. Do the descriptive statistics and do null value condition check, write an inference on it. 2.2 Data Split: Split the data into test and train, build classification model CART, Random Forest, Artificial Neural Network 2.3 Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC_AUC score for each model 2.4 Final Model: Compare all the model and write an inference which model is best/optimized. 2.5 Inference: Based on the whole Analysis, what are the business insights and recommendations

Dataset for Problem 2: insurance_part2_data-1.csv

Attribute Information:

1. Target: Claim Status (Claimed)
2. Code of tour firm (Agency_Code)
3. Type of tour insurance firms (Type)
4. Distribution channel of tour insurance agencies (Channel)
5. Name of the tour insurance products (Product)
6. Duration of the tour (Duration)
7. Destination of the tour (Destination)
8. Amount of sales of tour insurance policies (Sales)
9. The commission received for tour insurance firm (Commission)
10. Age of insured (Age)

2.1 Data Ingestion: Read the dataset. Do the descriptive statistics and do null value condition check, write an inference on it.

In [1]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, roc_curve, classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

Loading the data

In [2]:

```
dfi = pd.read_csv("insurance_part2_data.csv")
```

In [3]:

```
dfi.head()
```

Out[3]:

| Age | Agency_Code | Type | Claimed | Commision | Channel | Duration | Sales | Product Name | Destination |
|-----|-------------|------|---------|-----------|---------|----------|-------|--------------|-------------|
|-----|-------------|------|---------|-----------|---------|----------|-------|--------------|-------------|

| 0 | Age | Agency_Code | Type | Claimed | Commision | Channel | Duration | Sales | Product Name | Destination |
|---|-----|-------------|---------------|---------|-----------|---------|----------|-------|-------------------|-------------|
| 1 | 36 | EPX | Travel Agency | No | 0.00 | Online | 34 | 20.00 | Customised Plan | ASIA |
| 2 | 39 | CWT | Travel Agency | No | 5.94 | Online | 3 | 9.90 | Customised Plan | Americas |
| 3 | 36 | EPX | Travel Agency | No | 0.00 | Online | 4 | 26.00 | Cancellation Plan | ASIA |
| 4 | 33 | JZI | Airlines | No | 6.30 | Online | 53 | 18.00 | Bronze Plan | ASIA |

In [4]:

```
dfi.shape
```

Out[4]:

(3000, 10)

In [5]:

```
dfi.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Age                   3000 non-null   int64
1   Agency_Code           3000 non-null   object
2   Type                  3000 non-null   object
3   Claimed               3000 non-null   object
4   Commision             3000 non-null   float64
5   Channel               3000 non-null   object
6   Duration              3000 non-null   int64
7   Sales                 3000 non-null   float64
8   Product Name          3000 non-null   object
9   Destination           3000 non-null   object
dtypes: float64(2), int64(2), object(6)
memory usage: 234.5+ KB
```

In [6]:

```
dfi.describe()
```

Out[6]:

| | Age | Commision | Duration | Sales |
|-------|-------------|-------------|-------------|-------------|
| count | 3000.000000 | 3000.000000 | 3000.000000 | 3000.000000 |
| mean | 38.091000 | 14.529203 | 70.001333 | 60.249913 |
| std | 10.463518 | 25.481455 | 134.053313 | 70.733954 |
| min | 8.000000 | 0.000000 | -1.000000 | 0.000000 |
| 25% | 32.000000 | 0.000000 | 11.000000 | 20.000000 |
| 50% | 36.000000 | 4.630000 | 26.500000 | 33.000000 |
| 75% | 42.000000 | 17.235000 | 63.000000 | 69.000000 |
| max | 84.000000 | 210.210000 | 4580.000000 | 539.000000 |

In [7]:

```
dfi.isnull().sum()
```

Out[7]:

```
Age                0
Agency_Code       0
Type               0
Claimed            0
Commision          0
Channel            0
```

Duration 0
Sales 0
Product Name 0
Destination 0
dtype: int64

In [8]:

```
dfi.describe(include="all")
```

Out[8]:

| | Age | Agency_Code | Type | Claimed | Commision | Channel | Duration | Sales | Product Name | Destination |
|--------|-------------|-------------|---------------|---------|-------------|---------|-------------|-------------|-----------------|-------------|
| count | 3000.000000 | 3000 | 3000 | 3000 | 3000.000000 | 3000 | 3000.000000 | 3000.000000 | 3000 | 3000 |
| unique | NaN | 4 | 2 | 2 | NaN | 2 | NaN | NaN | 5 | 5 |
| top | NaN | EPX | Travel Agency | No | NaN | Online | NaN | NaN | Customised Plan | AS |
| freq | NaN | 1365 | 1837 | 2076 | NaN | 2954 | NaN | NaN | 1136 | 24 |
| mean | 38.091000 | NaN | NaN | NaN | 14.529203 | NaN | 70.001333 | 60.249913 | NaN | NaN |
| std | 10.463518 | NaN | NaN | NaN | 25.481455 | NaN | 134.053313 | 70.733954 | NaN | NaN |
| min | 8.000000 | NaN | NaN | NaN | 0.000000 | NaN | -1.000000 | 0.000000 | NaN | NaN |
| 25% | 32.000000 | NaN | NaN | NaN | 0.000000 | NaN | 11.000000 | 20.000000 | NaN | NaN |
| 50% | 36.000000 | NaN | NaN | NaN | 4.630000 | NaN | 26.500000 | 33.000000 | NaN | NaN |
| 75% | 42.000000 | NaN | NaN | NaN | 17.235000 | NaN | 63.000000 | 69.000000 | NaN | NaN |
| max | 84.000000 | NaN | NaN | NaN | 210.210000 | NaN | 4580.000000 | 539.000000 | NaN | NaN |

Getting unique counts of all nominal variables

In [9]:

```
dfi.columns
```

Out[9]:

Index(['Age', 'Agency_Code', 'Type', 'Claimed', 'Commision', 'Channel',
 'Duration', 'Sales', 'Product Name', 'Destination'],
 dtype='object')

In [10]:

```
for column in dfi[['Agency_Code', 'Type', 'Claimed', 'Channel',  
                  'Product Name', 'Destination']]:  
    print(column.upper(),': ',dfi[column].unique())  
    print(dfi[column].value_counts().sort_values())  
    print('\n')
```

AGENCY_CODE : 4
JZI 239
CWT 472
C2B 924
EPX 1365
Name: Agency_Code, dtype: int64

TYPE : 2
Airlines 1163
Travel Agency 1837
Name: Type, dtype: int64

CLAIMED : 2
Yes 924

```
No          2076
Name: Claimed, dtype: int64
```

```
CHANNEL :    2
Offline      46
Online     2954
Name: Channel, dtype: int64
```

```
PRODUCT NAME :    5
Gold Plan           109
Silver Plan         427
Bronze Plan         650
Cancellation Plan    678
Customised Plan     1136
Name: Product Name, dtype: int64
```

```
DESTINATION :    3
EUROPE        215
Americas      320
ASIA          2465
Name: Destination, dtype: int64
```

```
In [11]:
```

```
dups=dfi.duplicated()
dups
dups.sum()
```

```
Out[11]:
```

```
139
```

```
In [12]:
```

```
dfi.drop_duplicates(inplace=True)
```

```
In [13]:
```

```
dups=dfi.duplicated()
dups
dups.sum()
```

```
Out[13]:
```

```
0
```

```
In [14]:
```

```
dfi.shape
```

```
Out[14]:
```


```
(2861, 10)
```


```
In [15]:
```

```
plt.figure(figsize=(8,8))
dfi[['Age', 'Commision', 'Duration', 'Sales']].boxplot(vert=0)
```

```
Out[15]:
```

```
<AxesSubplot:>
```



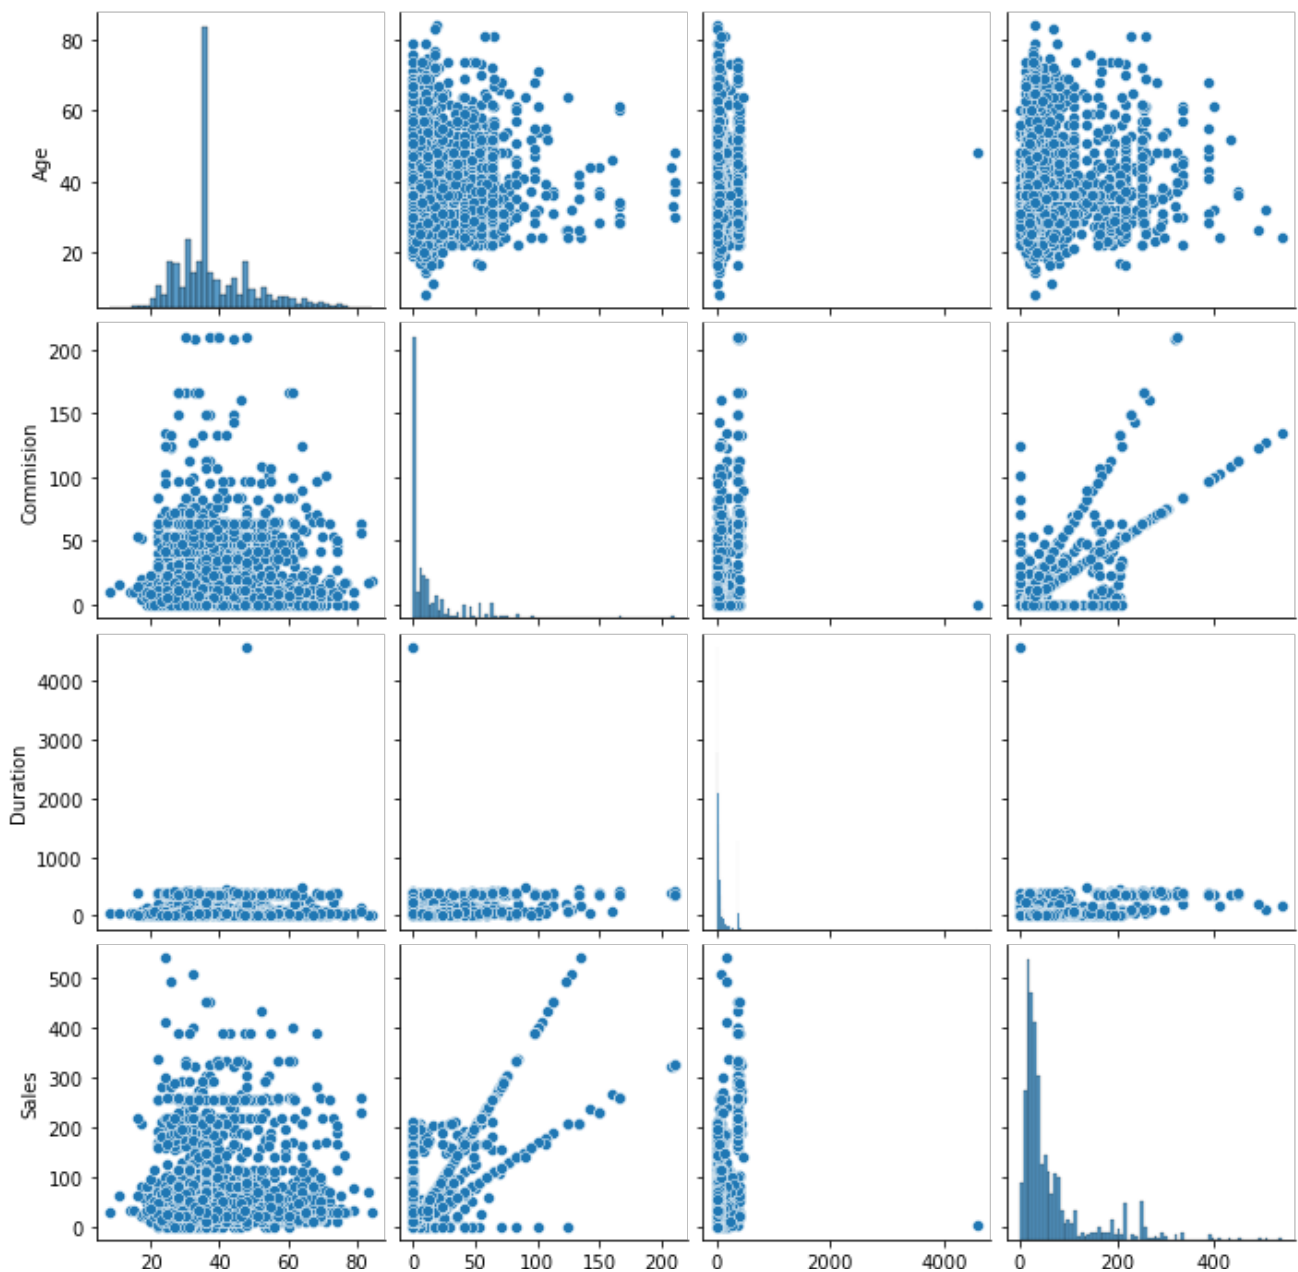


In [16]:

```
sns.pairplot(dfi)
```

Out[16]:

<seaborn.axisgrid.PairGrid at 0x1b041442c40>



Age

Commision

Duration

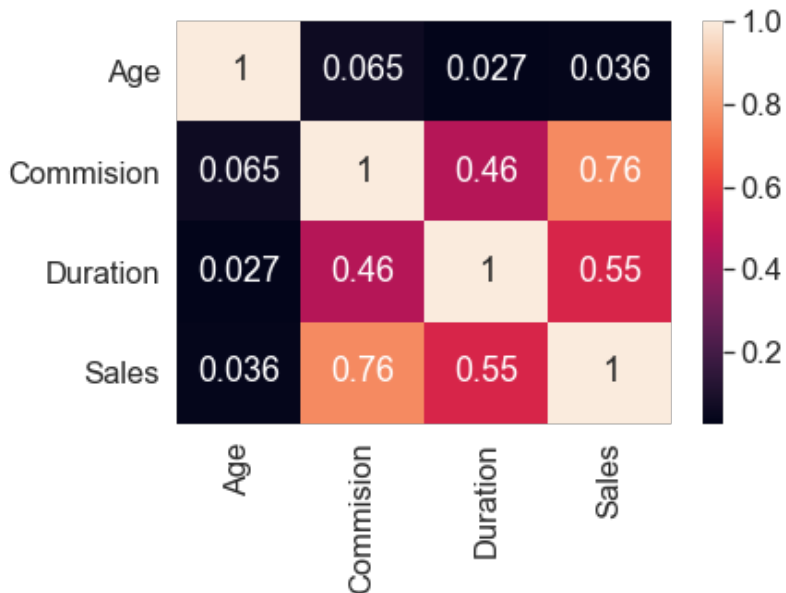
Sales

In [17]:

```
plt.figure(figsize=(6,4))
sns.set(font_scale=1.5)
sns.heatmap((dfi).corr(), annot=True)
```

Out[17]:

<AxesSubplot:>



In [18]:

```
for feature in dfi.columns:
    if dfi[feature].dtype == 'object':
        print('\n')
        print('feature:', feature)
        print(pd.Categorical(dfi[feature].unique()))
        print(pd.Categorical(dfi[feature].unique()).codes)
        dfi[feature] = pd.Categorical(dfi[feature]).codes
```

```
feature: Agency_Code
['C2B', 'EPX', 'CWT', 'JZI']
Categories (4, object): ['C2B', 'CWT', 'EPX', 'JZI']
[0 2 1 3]
```

```
feature: Type
['Airlines', 'Travel Agency']
Categories (2, object): ['Airlines', 'Travel Agency']
[0 1]
```

```
feature: Claimed
['No', 'Yes']
Categories (2, object): ['No', 'Yes']
[0 1]
```

```
feature: Channel
['Online', 'Offline']
Categories (2, object): ['Offline', 'Online']
[1 0]
```

```
feature: Product Name
['Customised Plan', 'Cancellation Plan', 'Bronze Plan', 'Silver Plan', 'Gold Plan']
Categories (5, object): ['Bronze Plan', 'Cancellation Plan', 'Customised Plan', 'Gold Plan', 'Silver Plan']
[2 1 0 4 3]
```

```
feature: Destination
['ASIA', 'Americas', 'EUROPE']
Categories (3, object): ['ASIA', 'Americas', 'EUROPE']
[0 1 2]
```

In [19]:

```
dfi.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 2861 entries, 0 to 2999
Data columns (total 10 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age             2861 non-null   int64
1   Agency_Code     2861 non-null   int8
2   Type            2861 non-null   int8
3   Claimed         2861 non-null   int8
4   Commision       2861 non-null   float64
5   Channel         2861 non-null   int8
6   Duration        2861 non-null   int64
7   Sales           2861 non-null   float64
8   Product Name    2861 non-null   int8
9   Destination     2861 non-null   int8
dtypes: float64(2), int64(2), int8(6)
memory usage: 128.5 KB
```

In [20]:

```
dfi.head()
```

Out[20]:

| | Age | Agency_Code | Type | Claimed | Commision | Channel | Duration | Sales | Product Name | Destination |
|---|-----|-------------|------|---------|-----------|---------|----------|-------|--------------|-------------|
| 0 | 48 | 0 | 0 | 0 | 0.70 | 1 | 7 | 2.51 | 2 | 0 |
| 1 | 36 | 2 | 1 | 0 | 0.00 | 1 | 34 | 20.00 | 2 | 0 |
| 2 | 39 | 1 | 1 | 0 | 5.94 | 1 | 3 | 9.90 | 2 | 1 |
| 3 | 36 | 2 | 1 | 0 | 0.00 | 1 | 4 | 26.00 | 1 | 0 |
| 4 | 33 | 3 | 0 | 0 | 6.30 | 1 | 53 | 18.00 | 0 | 0 |

proportion of 1s and 0s

In [21]:

```
dfi.Claimed.value_counts(normalize=True)
```

Out[21]:

```
0    0.680531
1    0.319469
Name: Claimed, dtype: float64
```

2.2 Data Split: Split the data into test and train, build classification model CART, Random Forest, Artificial Neural Network

Extracting the target column into separate vectors for training set and test set

In [22]:

```
X=dfi.drop("Claimed", axis=1)
Y=dfi.pop("Claimed")
X.head()
```

Out[22]:

| | Age | Agency_Code | Type | Commision | Channel | Duration | Sales | Product Name | Destination | |
|---|-----|-------------|------|-----------|---------|----------|-------|--------------|-------------|---|
| 0 | 48 | | 0 | 0 | 0.70 | 1 | 7 | 2.51 | 2 | 0 |
| 1 | 36 | | 2 | 1 | 0.00 | 1 | 34 | 20.00 | 2 | 0 |
| 2 | 39 | | 1 | 1 | 5.94 | 1 | 3 | 9.90 | 2 | 1 |
| 3 | 36 | | 2 | 1 | 0.00 | 1 | 4 | 26.00 | 1 | 0 |
| 4 | 33 | | 3 | 0 | 6.30 | 1 | 53 | 18.00 | 0 | 0 |

Splitting data into test and train

In [23]:

```
from sklearn.model_selection import train_test_split

X_train, X_test, train_labels, test_labels = train_test_split(X, Y, test_size=.30, random_state=1)
```

Checking the dimensions of the test and train data

In [24]:

```
print('X_test', X_test.shape)
print('X_train', X_train.shape)
print('test_labels', test_labels.shape)
print('train_labels', train_labels.shape)
```

```
X_test (859, 9)
X_train (2002, 9)
test_labels (859,)
train_labels (2002,)
```

Building a Decision Tree Classifier

In [25]:

```
param_grid = {
    'criterion': ['gini'],
    'max_depth': [5, 6, 7, 8, 9, 10],
    'min_samples_leaf': [25, 50, 75, 100, 125, 150],
    'min_samples_split': [75, 150, 225, 300, 375, 450],
}

dtcl = DecisionTreeClassifier(random_state=1)

grid_search = GridSearchCV(estimator = dtcl, param_grid = param_grid, cv = 10)
```

In [26]:

```
grid_search.fit(X_train, train_labels)
print(grid_search.best_params_)
best_grid = grid_search.best_estimator_
best_grid
#{'criterion': 'gini', 'max_depth': 10, 'min_samples_leaf': 50, 'min_samples_split': 450}

{'criterion': 'gini', 'max_depth': 5, 'min_samples_leaf': 50, 'min_samples_split': 300}
```

Out[26]:

```
DecisionTreeClassifier(max_depth=5, min_samples_leaf=50, min_samples_split=300,
                      random_state=1)
```

Generating Tree

In [27]:

```
train_char_label = ['no', 'yes']
tree_regularized = open('tree_regularized.dot', 'w')
dot_data = tree.export_graphviz(best_grid, out_file= tree_regularized , feature_names = list(X_train), class_names = list(train_char_label))

tree_regularized.close()
dot_data
```

<http://webgraphviz.com/>

Variable Importance

In [28]:

```
print (pd.DataFrame(best_grid.feature_importances_, columns = ["Imp"], index = X_train.columns).sort_values('Imp',ascending=False))
```

| | Imp |
|--------------|----------|
| Agency_Code | 0.601720 |
| Sales | 0.305611 |
| Product Name | 0.047457 |
| Duration | 0.016689 |
| Commision | 0.014763 |
| Age | 0.013760 |
| Type | 0.000000 |
| Channel | 0.000000 |
| Destination | 0.000000 |

Predicting on Training and Test dataset

In [29]:

```
ytrain_predict = best_grid.predict(X_train)
ytest_predict = best_grid.predict(X_test)
```

Getting the Predicted Classes and Probs

In [30]:

```
ytest_predict
ytest_predict_proba=best_grid.predict_proba(X_test)
ytest_predict_proba
pd.DataFrame(ytest_predict_proba).head()
```

Out[30]:

| | 0 | 1 |
|---|----------|----------|
| 0 | 0.573171 | 0.426829 |
| 1 | 0.971223 | 0.028777 |
| 2 | 0.232975 | 0.767025 |
| 3 | 0.837500 | 0.162500 |
| 4 | 0.837500 | 0.162500 |

Model Evaluation

AUC and ROC for the training data


In [31]:

```
# predict probabilities
probs = best_grid.predict_proba(X_train)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
cart_train_auc = roc_auc_score(train_labels, probs)
print('AUC: %.3f' % cart_train_auc)
# calculate roc curve
cart_train_fpr, cart_train_tpr, cart_train_thresholds = roc_curve(train_labels, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(cart_train_fpr, cart_train_tpr)
```

AUC: 0.809

Out[31]:

[<matplotlib.lines.Line2D at 0x1b0438bfeb0>]



AUC and ROC for the test data


In [32]:


```
# predict probabilities
probs = best_grid.predict_proba(X_test)
# keep probabilities for the positive outcome only
probs = probs[:, 1]
# calculate AUC
cart_test_auc = roc_auc_score(test_labels, probs)
print('AUC: %.3f' % cart_test_auc)
# calculate roc curve
cart_test_fpr, cart_test_tpr, cart_test_thresholds = roc_curve(test_labels, probs)
plt.plot([0, 1], [0, 1], linestyle='--')
# plot the roc curve for the model
plt.plot(cart_test_fpr, cart_test_tpr)
```

AUC: 0.796

Out[32]:

[<matplotlib.lines.Line2D at 0x1b043a2d610>]





Confusion Matrix for the training data

In [33]:

```
confusion_matrix(train_labels, ytrain_predict)
```

Out[33]:

```
array([[1157,  202],
       [ 270,  373]], dtype=int64)
```

In [34]:

```
#Train Data Accuracy
cart_train_acc=best_grid.score(X_train,train_labels)
cart_train_acc
```

Out[34]:

0.7642357642357642

In [35]:

```
print(classification_report(train_labels, ytrain_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.85 | 0.83 | 1359 |
| 1 | 0.65 | 0.58 | 0.61 | 643 |
| accuracy | | | 0.76 | 2002 |
| macro avg | 0.73 | 0.72 | 0.72 | 2002 |
| weighted avg | 0.76 | 0.76 | 0.76 | 2002 |

In [36]:

```
cart_metrics=classification_report(train_labels, ytrain_predict,output_dict=True)
df=pd.DataFrame(cart_metrics).transpose()
cart_train_f1=round(df.loc["1"][2],2)
cart_train_recall=round(df.loc["1"][1],2)
cart_train_precision=round(df.loc["1"][0],2)
print ('cart_train_precision ',cart_train_precision)
print ('cart_train_recall ',cart_train_recall)
print ('cart_train_f1 ',cart_train_f1)
```

```
cart_train_precision  0.65
cart_train_recall    0.58
cart_train_f1        0.61
```

Confusion Matrix for test data

In [37]:

```
confusion_matrix(test_labels, ytest_predict)
```

Out[37]:

```
array([[510,  78],
       [109, 162]], dtype=int64)
```

In [38]:

```
#Test Data Accuracy
cart_test_acc=best_grid.score(X_test,test_labels)
cart_test_acc
```

Out[38]:

0.7823050058207218

In [39]:

```
print(classification_report(test_labels, ytest_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.87 | 0.85 | 588 |
| 1 | 0.68 | 0.60 | 0.63 | 271 |
| accuracy | | | 0.78 | 859 |
| macro avg | 0.75 | 0.73 | 0.74 | 859 |
| weighted avg | 0.78 | 0.78 | 0.78 | 859 |

In [40]:

```
cart_metrics=classification_report(test_labels, ytest_predict,output_dict=True)
df=pd.DataFrame(cart_metrics).transpose()
cart_test_precision=round(df.loc["1"][0],2)
cart_test_recall=round(df.loc["1"][1],2)
cart_test_f1=round(df.loc["1"][2],2)
print ('cart_test_precision ',cart_test_precision)
print ('cart_test_recall ',cart_test_recall)
print ('cart_test_f1 ',cart_test_f1)
```

```
cart_test_precision  0.68
cart_test_recall    0.6
cart_test_f1        0.63
```

Cart Conclusion

CART CONCLUSION Train Data: AUC: 81% Accuracy: 76% Precision: 65% f1-Score: 61% recall: 58%

Test Data: AUC: 79.6% Accuracy: 78.2% Precision: 68% f1-Score: 63% recall: 60%

Training and Test set results somewhat closer, and with the overall measures high, the model is a good model.

Agency_Code is the most important variable for Claimed.

Building a Random Forest Classifier

Grid Search for finding out the optimal values for the hyper parameters

In [41]:

```
param_grid = {
    'max_depth': [10], ## 20,30,40
    'max_features': [6], ## 7,8,9
    'min_samples_leaf': [10], ## 50,100
    'min_samples_split': [50], ## 60,70
    'n_estimators': [300] ## 100,200
}

rfcl = RandomForestClassifier(random_state=1)

grid_search = GridSearchCV(estimator = rfcl, param_grid = param_grid, cv = 5)
```

In [42]:

```
grid_search.fit(X_train, train_labels)
```

Out[42]:

```
GridSearchCV(cv=5, estimator=RandomForestClassifier(random_state=1),
             param_grid={'max_depth': [10], 'max_features': [6],
                          'min_samples_leaf': [10], 'min_samples_split': [50],
                          'n_estimators': [300]})
```

In [43]:

```
grid_search.best_params_
```

Out[43]:

```
{'max_depth': 10,
 'max_features': 6,
 'min_samples_leaf': 10,
 'min_samples_split': 50,
 'n_estimators': 300}
```

In [44]:

```
best_grid = grid_search.best_estimator_
```

In [45]:

```
best_grid
```

Out[45]:

```
RandomForestClassifier(max_depth=10, max_features=6, min_samples_leaf=10,
                       min_samples_split=50, n_estimators=300, random_state=1)
```

Predicting the Training and Testing data

In [46]:

```
ytrain_predict = best_grid.predict(X_train)
ytest_predict = best_grid.predict(X_test)
```

RF Model Performance Evaluation on Training data

In [47]:

```
confusion_matrix(train_labels, ytrain_predict)
```

Out[47]:

```
array([[1228, 131],
       [ 258, 385]], dtype=int64)
```

In [48]:

```
rf_train_acc=best_grid.score(X_train, train_labels)
rf_train_acc
```

Out[48]:

```
0.8056943056943057
```

In [49]:

```
print(classification_report(train_labels, ytrain_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.83 | 0.90 | 0.86 | 1359 |
| 1 | 0.75 | 0.60 | 0.66 | 643 |
| accuracy | | | 0.81 | 2002 |
| macro avg | 0.79 | 0.75 | 0.76 | 2002 |
| weighted avg | 0.80 | 0.81 | 0.80 | 2002 |

In [50]:


```
rf_metrics=classification_report(train_labels, ytrain_predict,output_dict=True)
df=pd.DataFrame(rf_metrics).transpose()
rf_train_precision=round(df.loc["1"][0],2)
rf_train_recall=round(df.loc["1"][1],2)
rf_train_f1=round(df.loc["1"][2],2)
print ('rf_train_precision ',rf_train_precision)
print ('rf_train_recall ',rf_train_recall)
print ('rf_train_f1 ',rf_train_f1)
```

```
rf_train_precision  0.75
rf_train_recall    0.6
rf_train_f1        0.66
```

In [51]:

```
rf_train_fpr, rf_train_tpr, _=roc_curve(train_labels,best_grid.predict_proba(X_train)[: ,1
])
plt.plot(rf_train_fpr,rf_train_tpr,color='green')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
rf_train_auc=roc_auc_score(train_labels,best_grid.predict_proba(X_train)[: ,1])
print('Area under Curve is', rf_train_auc)
```

Area under Curve is 0.8639002468423744



RF Model Performance Evaluation on Test data

In [52]:

```
confusion_matrix(test_labels,ytest_predict)
```

Out[52]:

```
array([[522,  66],
       [118, 153]], dtype=int64)
```

In [53]:

```
rf_test_acc=best_grid.score(X_test,test_labels)
rf_test_acc
```

Out[53]:

```
0.7857974388824214
```

In [54]:

```
print(classification_report(test_labels,ytest_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.82 | 0.89 | 0.85 | 588 |
| 1 | 0.70 | 0.56 | 0.62 | 271 |
| accuracy | | | 0.79 | 859 |
| macro avg | 0.76 | 0.73 | 0.74 | 859 |
| weighted avg | 0.78 | 0.79 | 0.78 | 859 |

In [55]:

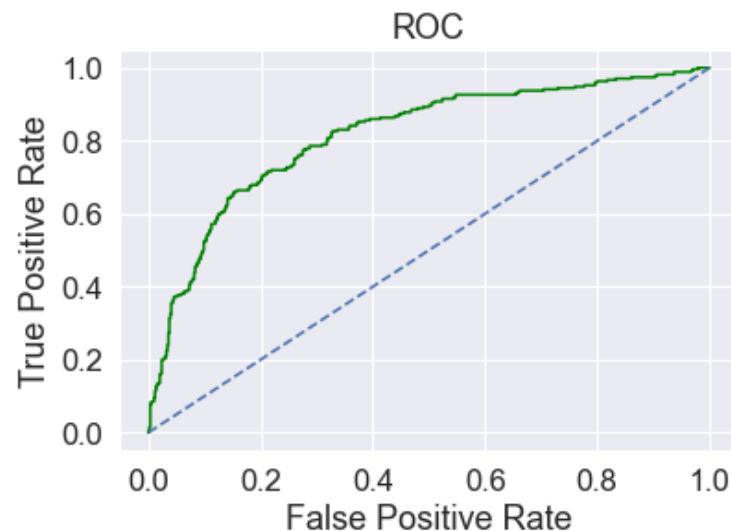
```
rf_metrics=classification_report(test_labels, ytest_predict,output_dict=True)
df=pd.DataFrame(rf_metrics).transpose()
rf_test_precision=round(df.loc["1"][0],2)
rf_test_recall=round(df.loc["1"][1],2)
rf_test_f1=round(df.loc["1"][2],2)
print ('rf_test_precision ',rf_test_precision)
print ('rf_test_recall ',rf_test_recall)
print ('rf_test_f1 ',rf_test_f1)
```

```
rf_test_precision 0.7
rf_test_recall 0.56
rf_test_f1 0.62
```

In [56]:

```
rf_test_fpr, rf_test_tpr, _=roc_curve(test_labels,best_grid.predict_proba(X_test)[:,1])
plt.plot(rf_test_fpr,rf_test_tpr,color='green')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
rf_test_auc=roc_auc_score(test_labels,best_grid.predict_proba(X_test)[:,1])
print('Area under Curve is', rf_test_auc)
```

Area under Curve is 0.8136186208800863



In [57]:

```
# Variable Importance
print (pd.DataFrame(best_grid.feature_importances_, columns = ["Imp"], index = X_train.columns).sort_values('Imp',ascending=False))
```

| | Imp |
|--------------|----------|
| Agency_Code | 0.325075 |
| Sales | 0.207197 |
| Product Name | 0.169962 |
| Duration | 0.101869 |
| Commision | 0.095711 |
| Age | 0.069549 |
| Tvne | 0.015744 |

```
Destination    0.013474
Channel        0.001419
```

Random Forest Conclusion

RF conclusion Train Data: AUC: 86.3% Accuracy: 80.6% Precision: 75% f1-Score: 66% Recall: 60%

Test Data: AUC: 81.3% Accuracy: 78.6% Precision: 70% f1-Score: 62% Recall: 56% Training and Test set results are almost similar, and with the overall measures high, the model is a good model. Agency_Code is again the most important variable for predicting claimed

Building a Neural Network Classifier

In [58]:

```
from sklearn.preprocessing import StandardScaler
```

In [59]:

```
#Initialize an object for StandardScaler
sc = StandardScaler()
```

In [60]:

```
#Scale the training data
X_train = sc.fit_transform(X_train)
```

In [61]:

```
X_train
```

Out[61]:

```
array([[ 2.88764239, -1.2626112 , -1.19813318, ..., -0.65375471,
        -1.31338076, -0.44775345],
       [-0.21666128,  0.71683095,  0.83463176, ..., -0.37032806,
         0.24339146, -0.44775345],
       [ 2.04101412, -0.27289013,  0.83463176, ...,  0.11574864,
         0.24339146,  1.24676906],
       ...,
       [-0.21666128,  0.71683095,  0.83463176, ..., -0.68209737,
        -0.53499465, -0.44775345],
       [-0.21666128,  0.71683095,  0.83463176, ...,  0.72086453,
         0.24339146, -0.44775345],
       [-0.21666128,  0.71683095,  0.83463176, ...,  0.72086453,
         0.24339146,  1.24676906]])
```

In [62]:

```
# Apply the transformation on the test data
X_test = sc.transform(X_test)
```

In [63]:

```
X_test
```

Out[63]:

```
array([[ -0.68701032, -0.27289013,  0.83463176, ...,  0.50829455,
         0.24339146, -0.44775345],
       [  2.79357258,  0.71683095,  0.83463176, ..., -0.45535606,
        -0.53499465, -0.44775345],
       [  0.34775757, -1.2626112 , -1.19813318, ...,  0.32406723,
         1.80016368, -0.44775345],
       ...,
       [  1.19438584, -1.2626112 , -1.19813318, ..., -0.63958338,
        -1.31338076, -0.44775345],
```



```
[ 1.38252546,  0.71683095,  0.83463176, ..., -0.56872671,
 0.24339146, -0.44775345],
[-0.21666128,  0.71683095,  0.83463176, ..., -0.56872671,
 0.24339146, -0.44775345]])
```

In [64]:

```
param_grid = {
    'hidden_layer_sizes': [100], # 50, 200
    'max_iter': [2500], #5000,2500
    'solver': ['adam'], #sgd
    'tol': [0.01],
}

nncl = MLPClassifier(random_state=1)

grid_search = GridSearchCV(estimator = nncl, param_grid = param_grid, cv = 10)
```

In [65]:

```
grid_search.fit(X_train, train_labels)
grid_search.best_params_
#{'hidden_layer_sizes': 100, 'max_iter': 2500, 'solver': 'adam', 'tol': 0.01}
```

Out[65]:

```
{'hidden_layer_sizes': 100, 'max_iter': 2500, 'solver': 'adam', 'tol': 0.01}
```

In [66]:

```
best_grid = grid_search.best_estimator_
best_grid
```

Out[66]:

```
MLPClassifier(hidden_layer_sizes=100, max_iter=2500, random_state=1, tol=0.01)
```

Predicting the Training and Testing data

In [67]:

```
ytrain_predict = best_grid.predict(X_train)
ytest_predict = best_grid.predict(X_test)
```

NN Model Performance Evaluation on Training data

In [68]:

```
confusion_matrix(train_labels,ytrain_predict)
```

Out[68]:

```
array([[1159,  200],
       [ 285,  358]], dtype=int64)
```

In [69]:

```
nn_train_acc=best_grid.score(X_train,train_labels)
nn_train_acc
```

Out[69]:

```
0.7577422577422578
```

In [70]:

```
print(classification_report(train_labels,ytrain_predict))
```

```
precision    recall  f1-score   support
```

| | | | | |
|--------------|------|------|------|------|
| 0 | 0.80 | 0.85 | 0.83 | 1359 |
| 1 | 0.64 | 0.56 | 0.60 | 643 |
| accuracy | | | 0.76 | 2002 |
| macro avg | 0.72 | 0.70 | 0.71 | 2002 |
| weighted avg | 0.75 | 0.76 | 0.75 | 2002 |

In [71]:


```
nn_metrics=classification_report(train_labels, ytrain_predict,output_dict=True)
df=pd.DataFrame(nn_metrics).transpose()
nn_train_precision=round(df.loc["1"][0],2)
nn_train_recall=round(df.loc["1"][1],2)
nn_train_f1=round(df.loc["1"][2],2)
print ('nn_train_precision ',nn_train_precision)
print ('nn_train_recall ',nn_train_recall)
print ('nn_train_f1 ',nn_train_f1)
```

```
nn_train_precision  0.64
nn_train_recall    0.56
nn_train_f1        0.6
```

In [72]:

```
nn_train_fpr, nn_train_tpr, _=roc_curve(train_labels,best_grid.predict_proba(X_train)[:,1]
])
plt.plot(nn_train_fpr,nn_train_tpr,color='black')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
nn_train_auc=roc_auc_score(train_labels,best_grid.predict_proba(X_train)[:,1])
print('Area under Curve is', nn_train_auc)
```

Area under Curve is 0.7921265636497425



NN Model Performance Evaluation on Test data

In [73]:

```
confusion_matrix(test_labels,ytest_predict)
```

Out[73]:

```
array([[511,  77],
       [122, 149]], dtype=int64)
```

In [74]:

```
nn_test_acc=best_grid.score(X_test,test_labels)
```

```
nn_test_acc
```

```
Out[74]:
```

```
0.7683352735739232
```

```
In [75]:
```

```
print(classification_report(test_labels,ytest_predict))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.81 | 0.87 | 0.84 | 588 |
| 1 | 0.66 | 0.55 | 0.60 | 271 |
| accuracy | | | 0.77 | 859 |
| macro avg | 0.73 | 0.71 | 0.72 | 859 |
| weighted avg | 0.76 | 0.77 | 0.76 | 859 |

```
In [76]:
```


```
nn_metrics=classification_report(test_labels, ytest_predict,output_dict=True)
df=pd.DataFrame(nn_metrics).transpose()
nn_test_precision=round(df.loc["1"][0],2)
nn_test_recall=round(df.loc["1"][1],2)
nn_test_f1=round(df.loc["1"][2],2)
print ('nn_test_precision ',nn_test_precision)
print ('nn_test_recall ',nn_test_recall)
print ('nn_test_f1 ',nn_test_f1)
```

```
nn_test_precision 0.66
nn_test_recall 0.55
nn_test_f1 0.6
```

```
In [77]:
```

```
nn_test_fpr, nn_test_tpr, _=roc_curve(test_labels,best_grid.predict_proba(X_test)[: ,1])
plt.plot(nn_test_fpr,nn_test_tpr,color='black')
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
nn_test_auc=roc_auc_score(test_labels,best_grid.predict_proba(X_test)[: ,1])
print('Area under Curve is', nn_test_auc)
```

```
Area under Curve is 0.7977947636619223
```



```
In [78]:
```

```
best_grid.score
```

```
Out[78]:
```

```
0.7683352735739232
```

```
<bound method ClassifierMixin.score of MLPClassifier(hidden_layer_sizes=100, max_iter=2500, random_state=1, tol=0.01)>
```

```
In [ ]:
```

Neural Network Conclusion

Train Data:

AUC: 79.15% Accuracy: 75.77% Precision: 64% f1-Score: 60% Recall : 56%

Test Data: AUC: 79.8% Accuracy: 76.8% Precision: 66% f1-Score: 60% Recall : 55% Training and Test set results are almost similar, and with the overall measures high, the model is a good model

Training and Test set results are almost similar, and with the overall measures high, the model is a good model.

Final Conclusion

Comparison of the performance metrics from the 3 models

```
In [79]:
```

```
index=['Accuracy', 'AUC', 'Recall', 'Precision', 'F1 Score']
data = pd.DataFrame({'CART Train':[cart_train_acc, cart_train_auc, cart_train_recall, cart_train_precision, cart_train_f1],
                    'CART Test':[cart_test_acc, cart_test_auc, cart_test_recall, cart_test_precision, cart_test_f1],
                    'Random Forest Train':[rf_train_acc, rf_train_auc, rf_train_recall, rf_train_precision, rf_train_f1],
                    'Random Forest Test':[rf_test_acc, rf_test_auc, rf_test_recall, rf_test_precision, rf_test_f1],
                    'Neural Network Train':[nn_train_acc, nn_train_auc, nn_train_recall, nn_train_precision, nn_train_f1],
                    'Neural Network Test':[nn_test_acc, nn_test_auc, nn_test_recall, nn_test_precision, nn_test_f1]}, index=index)
round(data, 2)
```

```
Out[79]:
```

| | CART Train | CART Test | Random Forest Train | Random Forest Test | Neural Network Train | Neural Network Test |
|-----------|------------|-----------|---------------------|--------------------|----------------------|---------------------|
| Accuracy | 0.76 | 0.78 | 0.81 | 0.79 | 0.76 | 0.77 |
| AUC | 0.81 | 0.80 | 0.86 | 0.81 | 0.79 | 0.80 |
| Recall | 0.58 | 0.60 | 0.60 | 0.56 | 0.56 | 0.55 |
| Precision | 0.65 | 0.68 | 0.75 | 0.70 | 0.64 | 0.66 |
| F1 Score | 0.61 | 0.63 | 0.66 | 0.62 | 0.60 | 0.60 |

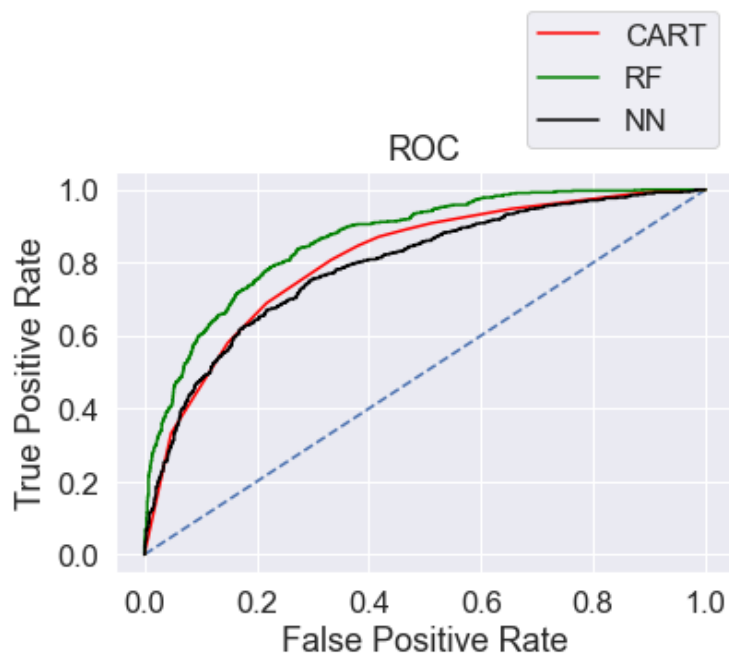
ROC Curve for the 3 models on the Training data

```
In [80]:
```

```
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(cart_train_fpr, cart_train_tpr, color='red', label="CART")
plt.plot(rf_train_fpr, rf_train_tpr, color='green', label="RF")
plt.plot(nn_train_fpr, nn_train_tpr, color='black', label="NN")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.legend(bbox_to_anchor=(0., 1.02, 1., .102), loc='lower right')
```

Out[80]:

<matplotlib.legend.Legend at 0x1b04395ab50>



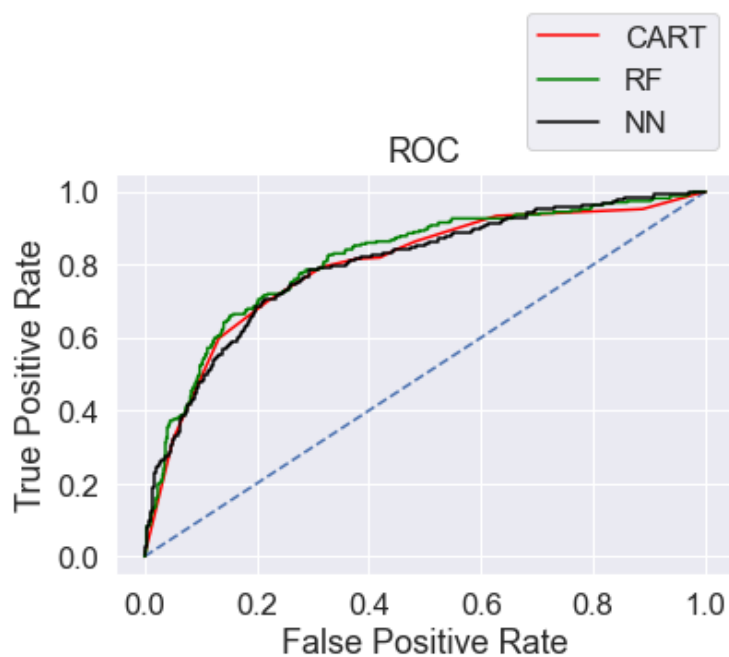
ROC Curve for the 3 models on the Test data

In [81]:

```
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(cart_test_fpr, cart_test_tpr, color='red', label="CART")
plt.plot(rf_test_fpr, rf_test_tpr, color='green', label="RF")
plt.plot(nn_test_fpr, nn_test_tpr, color='black', label="NN")
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.legend(bbox_to_anchor=(0., 1.02, 1., .102), loc='lower right')
```

Out[81]:

<matplotlib.legend.Legend at 0x1b04399dd00>



2.5 Inference: Based on the whole Analysis, what are the business insights and recommendations

Out of the 3 models, CART has slightly better performance than the RF and Neural network model based on

recall. Based on precision, we prefer RF over CART and ANN. Overall all the 3 models are reasonably stable enough to be used for making any future predictions. From Cart and Random Forest Model, the variable Agency code is found to be the most useful feature amongst all other features for predicting the target claimed .

```
In [ ]:
```

```
In [ ]:
```